

# **CREATING A WEBSITE USING FLASK+DYNAMODB+DOCKER**

**CLASS:**  
**INSTRUCTOR:**

## Introduction:

In this tutorial you will be learning how to create a simple website using flask web framework and integrate aws dynamodb database to it using python programming language.

- Flask

[Flask](#) is a micro web framework which is written in python which will help us to create web applications. At the most basic level it provides a development server and debugger. Flask depends on [jinja](#) template engine for template creation.

- DynamoDB

Amazon [DynamoDB](#) is a fully managed, serverless, key-value NoSQL database designed to run high-performance applications at any scale. DynamoDB offers built-in security, continuous backups, automated multi-Region replication, in-memory caching, and data export tools.

- Docker

We use docker container to package our website and deploy it in aws environment

## Prerequisites:

- Linux CLI (bash)
- Docker
- Python3 programming
- pip
- Flask
- AWS and IAM
- AWS elastic beanstalk cli

If you do not have docker installed in your machine please follow a guide to how to install docker on your machine. This is a [guide](#) on how you could install docker on a ubuntu machine.

You need python3 installed in your local machine if it is already there you need to download it from [here](#) and install it.

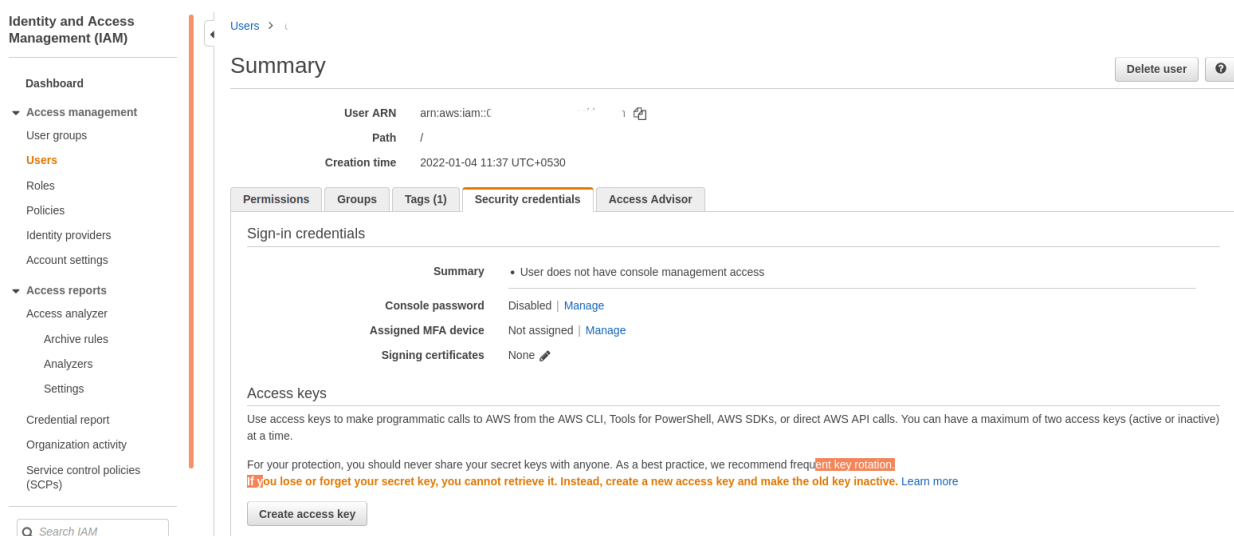
Pip will be available once you have installed python. After that you will need to install aws cli if you have not already configured the aws cli in your local machine. In a terminal try to paste below and install aws cli.

**pip install awscli --upgrade --user**

After that you need to configure the aws then paste below command and provide correct values which are related to your aws account

## aws configure

You will get below to add the data. For your user which is used to login to aws console must have an access key and secret. You have to remember the secret value you have entered when you are creating the access key. Please login to aws console and search for IAM and open IAM console navigate to your user and you will find the place to create a key if you have not already done it. Please follow this [guide](#) to get more details.



AWS Access Key ID [None]: \*\*\*PASTE ACCESS KEY ID HERE\*\*\*

AWS Secret Access Key [None]: \*\*\*PASTE SECRET ACCESS KEY HERE\*\*\*

Default region name [None]: \*\*\*TYPE YOUR PREFERRED REGION\*\*\*

Default output format [None]: json

Then we need to install elastic beanstalk. First clone the files from git using below command

git clone <https://github.com/aws/aws-elastic-beanstalk-cli-setup.git>

Then install eb using this command

```
python3.7 ./aws-elastic-beanstalk-cli-setup/scripts/ebcli_installer.py
```

After installation output will instruct you to export environment variable do as it mentioned. Command will look similar to this

```
echo 'export PATH="/home/xxxxx/.ebcli-virtual-env/executables:$PATH"' >>  
~/.bash_profile && source ~/.bash_profile
```

## **Project Setup:**

Ok well done. Now we are ready to go with our project setup. First let's create a directory to hold our project data. open up a terminal and create a directory and navigate to it.

```
~/Documents/tutorial$ mkdir flaskapp  
~/Documents/tutorial$ cd flaskapp/  
~/Documents/tutorial/flaskapp$
```

Now let's create a python virtual environment to install dependencies which we require to develop our app. In same terminal just paste below commands

```
~/Documents/tutorial/flaskapp$ python3 -m venv env  
~/Documents/tutorial/flaskapp$ source env/bin/activate
```

Now we could use pip command to install flask and boto3 libs for our project run below command in your terminal where you have activated the virtual environment. This is required for local testing later will add these dependencies in a requirements file for docker image creation.

```
pip install flask
```

```
pip install boto3
```

Ok, Then we need to instruct flask which file needs to be used as the main file so let's create the environment variable. Please add the code below in the terminal.

```
export FLASK_APP=application.py
```

Next create a directory named templates to keep our html files.

```
mkdir templates
```

Create a index.html file inside the templates directory.

```
~/Documents/tutorial/flaskapp$ mkdir templates
~/Documents/tutorial/flaskapp$ cd templates/
~/Documents/tutorial/flaskapp/templates$ touch index.html
```

Add below code in the index.html file.

```
<!DOCTYPE html>

<html lang="en">

<head>

    <meta charset="UTF-8">

    <title>FlaskDynamodb</title>

</head>

<body>

    <h1>Welcome to FlaskSite</h1>

</body>

</html>
```

Now let's go to the root directory and create an application.py file.

CMD: touch application.py

Try to add this code in application.py file

```
from flask import Flask, render_template
```

```
application = Flask(__name__)

@app.route('/')

def index():

    return render_template('index.html')
```

@app.route('/') indicate that if the server get an request for root (/) render the index.html file.

Ok now we can run our very first flask web app. In terminal run:

flask run

```
* Serving Flask app "app.py"
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)
```

Now navigate to the link <http://127.0.0.1:5000/> and you will see a web page as below.

A screenshot of a web browser window. The address bar shows '127.0.0.1:5000'. The page content displays 'Welcome to FlaskSite' in a bold, black font.

## Database creation:

Create a json file named db-table.json and add this content in it.

```
{

    "TableName": "FilmInfoTable",

    "KeySchema": [

        { "AttributeName": "Director", "KeyType": "HASH" },
```

```
{ "AttributeName": "FilmTitle", "KeyType": "RANGE" }

],

"AttributeDefinitions": [

  { "AttributeName": "Director", "AttributeType": "S" },

  { "AttributeName": "FilmTitle", "AttributeType": "S" }

],

"ProvisionedThroughput": {

  "ReadCapacityUnits": 5,

  "WriteCapacityUnits": 5

}

}
```

Here we are creating a dynamodb table named FilmInfoTable and it's partition key is Director and the sort key is the FilmTitle

### Partition key

The partition key is part of the table's primary key. It is a hash value that is used to retrieve items from your table and allocate data across hosts for scalability and availability.

### Sort key

You can use a sort key as the second part of a table's primary key. The sort key allows you to sort or search among all items sharing the same partition key.

More about the provisioned information can be found [here](#).

Let's create the table and add the data to the table now.

Run the below command in the terminal where your json file exists.

**aws dynamodb create-table --cli-input-json file://db-table.json**

You will see a output like this:

```
{
  "TableDescription": {
    "AttributeDefinitions": [
      {
        "AttributeName": "Director",
        "AttributeType": "S"
      },
      {
        "AttributeName": "FilmTitle",
        "AttributeType": "S"
      }
    ],
    "TableName": "FilmInfoTable",
    "KeySchema": [
      {
        "AttributeName": "Director",
        "KeyType": "HASH"
      },
      {
        "AttributeName": "FilmTitle",
        "KeyType": "RANGE"
      }
    ],
    "TableStatus": "CREATING",
    "CreationDateTime": 1648446951.308,
    "ProvisionedThroughput": {
      "NumberOfDecreasesToday": 0,
      "ReadCapacityUnits": 5,
      "WriteCapacityUnits": 5
    },
    "TableSizeBytes": 0,
    "ItemCount": 0,
    "TableArn": "arn:aws:dynamodb:eu-west-1:057186357501:table/FilmInfoTable",
    "TableId": "edcc9375-36f5-49ac-bd4c-7dde8e425649"
  }
}
```

You can navigate to aws console and search for the dynamodb and go dynamodb and see the table which you created.

The screenshot shows the AWS DynamoDB console interface. On the left is a sidebar with navigation options: Dashboard, Tables (selected), Update settings, Explore items, PartiQL editor, Backups, Exports to S3, and Reserved capacity. The main area displays 'DynamoDB > Tables'. At the top, there's a 'Tables (1)' header with an 'Info' tab, a search bar, and buttons for 'Actions', 'Delete', and 'Create table'. Below this is a table listing the tables. The table has columns: Name, Status, Partition key, Sort key, Indexes, Read capacity mode, Write capacity mode, Size, and Table class. One table is listed: 'FilmInfoTable' with status 'Active', partition key 'Director (S)', sort key 'FilmTitle (S)', 0 indexes, 'Provisioned (S)' for both read and write capacity, 0 bytes size, and 'DynamoDB Standard' table class.

<input type="checkbox"/>	Name	Status	Partition key	Sort key	Indexes	Read capacity mode	Write capacity mode	Size	Table class
<input type="checkbox"/>	FilmInfoTable	Active	Director (S)	FilmTitle (S)	0	Provisioned (S)	Provisioned (S)	0 bytes	DynamoDB Standard

Let's add some data to this table now. Create a json file named test-data.json and add the below content in it.

```
{

  "FilmInfoTable": [{

    "PutRequest": {
```



```
        "Item": {
            "Director": {
                "S": "ADAM"
            },
            "FilmTitle": {
                "S": "Sleep less"
            }
        }
    },
    {
        "PutRequest": {
            "Item": {
                "Director": {
                    "S": "Stefan"
                },
                "FilmTitle": {
                    "S": "Paint ball"
                }
            }
        }
    },
    {
```

```
{
  "PutRequest": {
    "Item": {
      "Director": {
        "S": "Lisa"
      },
      "FilmTitle": {
        "S": "Fedup"
      }
    }
  }
}
```

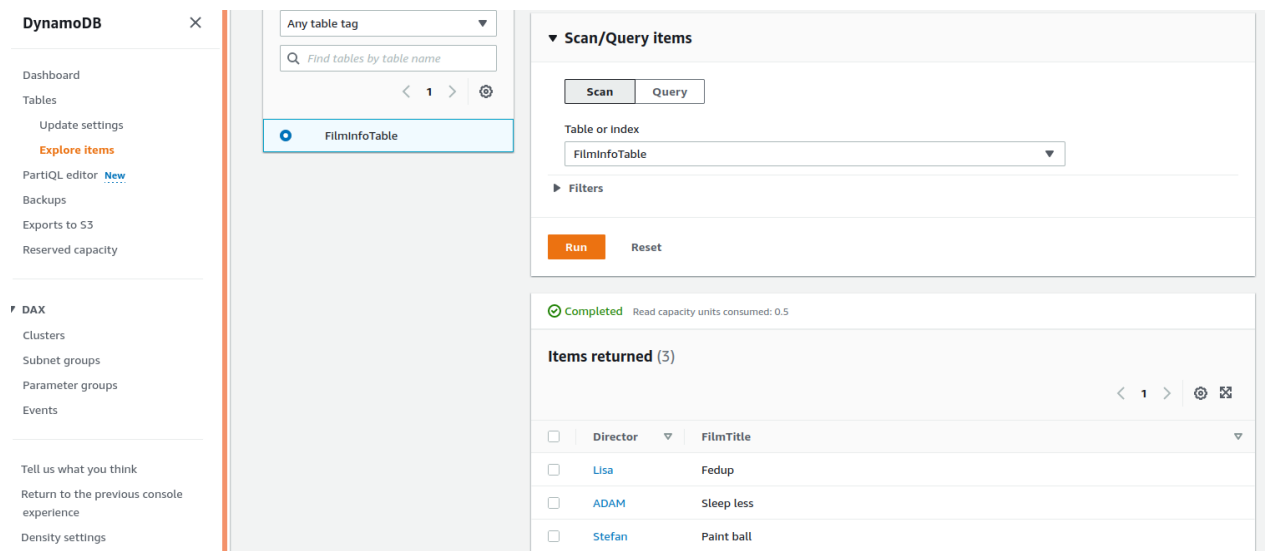
Now execute below command.

**aws dynamodb batch-write-item --request-items file:///test-data.json**

You should see a result like this

```
{"UnprocessedItems": {}}
```

If you explore the items in the newly created table you will see the inserted data.



The screenshot shows the AWS DynamoDB console interface. On the left is a navigation sidebar with options like Dashboard, Tables, Update settings, Explore items, PartiQL editor, Backups, Exports to S3, Reserved capacity, DAX, Clusters, Subnet groups, Parameter groups, and Events. The main panel is titled 'Scan/Query items' and shows a table named 'FilmInfoTable' selected. Below the table name, there are 'Run' and 'Reset' buttons. A status bar indicates 'Completed' with 'Read capacity units consumed: 0.5'. Below this, a table titled 'Items returned (3)' displays the data:

	Director	FilmTitle
<input type="checkbox"/>	Lisa	Fedup
<input type="checkbox"/>	ADAM	Sleep less
<input type="checkbox"/>	Stefan	Paint ball

Let's read these data from our flask code and visualize it in our web template. Create a file named `aws-controller.py` and add this code

```
import boto3

dynamo_client = boto3.client('dynamodb')

def get_items():

    response=dynamo_client.scan(

        TableName='FilmInfoTable'

    )

    return response['Items']
```

In this code we are using boto3 to connect our dynamodb database and fetch the table data.

After this we need to pass the fetched values to web page via app.py. So you need to change the app.py like this

```
from flask import Flask, render_template

import aws_controller

application = Flask(__name__)

@app.route('/')

def index():

    data=aws_controller.get_items()

    return render_template('index.html',films=data)
```

So here we are passing fetched film data to our template `render_template('index.html',films=data)` and now we need to read those data via our `index.html` file.

```
<!DOCTYPE html>

<html lang="en">

<head>

    <meta charset="UTF-8">

    <title>FlaskDynamodb</title>

</head>

<body>
```

```

<h1>Welcome to FlaskSite</h1>

{% for film in films %}

    <h2>{{ film['Director']['S'] }}</h2>

    <h3>{{ film['FilmTitle']['S'] }}</h3>

    <hr>

{% endfor %}

</body>

</html>

```

We use jinja template capabilities to loop “{% for film in films %}” through the film data we passed and visualize those. Run again the application and see the results.



Ok if you see the above result you have done a great job. Now let’s create a docker container out of this code and deploy it in AWS using aws Elastic Beanstalk.

## Docker Container Creation:

First let's create a requirements.txt file and add flask and boto3 in it.

```
flask
boto3
```

Create dockerfile in the root directory and add below content

```
# syntax=docker/dockerfile:1

FROM python:3.8-slim-buster

WORKDIR /python-docker

COPY requirements.txt requirements.txt
RUN pip3 install -r requirements.txt
ENV FLASK_APP=application.py

COPY . .
EXPOSE 80
EXPOSE 8000
CMD [ "python3", "-m" , "flask", "run", "--host=0.0.0.0", "--port=8000"]
```

We use python:3.8-slim-buster as the base image of our solution. Then we copy the requirement.txt file install requirements. We are copying the rest of the files using COPY .. and expose the port 8000 in the container. Then change the aws\_controller.py code to this.

Make sure you set AWS\_SERVER\_PUBLIC\_KEY and AWS\_SERVER\_SECRET\_KEY values to your values. It is not recommended to hard code these values but for learning purposes we hardcoded those.

```
from urllib import response
import boto3
import os

AWS_SERVER_PUBLIC_KEY="xxxxxxxxxxxxxxxx"
AWS_SERVER_SECRET_KEY="xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx"

dynamo_client = boto3.client('dynamodb',
                              aws_access_key_id=AWS_SERVER_PUBLIC_KEY,
                              aws_secret_access_key=AWS_SERVER_SECRET_KEY,
                              region_name='eu-west-1')
```

```
def get_items():
    response=dynamo_client.scan(
        TableName='FilmInfoTable'
    )
    return response['Items']
```

Here we are passing the credentials to access aws services via environment variables which we pass through the docker file.

Let's build our first docker file and run it locally first.

Open a terminal in the root folder and run Docker build command. Here we are tagging our image as mysite

```
docker build --tag mysite .
```

After you successfully build the image you could run the image as below. Here we are mapping port 5000 to exposed port 5000.

```
docker run -p 5000:5000 mysite
```

## Deploy using elastic beanstalk:

Open a new terminal and navigate to your root directory and run eb init.

eb init

It will ask to select a region, add the number of your region. In my case it is 4 because I'm using eu-west-1. It will ask for the app name, keep it as it is and select the python env as instruct finally type n for ssh access because we do not need to access the ec2 instance we create. See below image what I have entered.

```

Select a default region
1) us-east-1 : US East (N. Virginia)
2) us-west-1 : US West (N. California)
3) us-west-2 : US West (Oregon)
4) eu-west-1 : EU (Ireland)
5) eu-central-1 : EU (Frankfurt)
6) ap-south-1 : Asia Pacific (Mumbai)
7) ap-southeast-1 : Asia Pacific (Singapore)
8) ap-southeast-2 : Asia Pacific (Sydney)
9) ap-northeast-1 : Asia Pacific (Tokyo)
10) ap-northeast-2 : Asia Pacific (Seoul)
11) sa-east-1 : South America (Sao Paulo)
12) cn-north-1 : China (Beijing)
13) cn-northwest-1 : China (Ningxia)
14) us-east-2 : US East (Ohio)
15) ca-central-1 : Canada (Central)
16) eu-west-2 : EU (London)
17) eu-west-3 : EU (Paris)
18) eu-north-1 : EU (Stockholm)
19) eu-south-1 : EU (Milano)
20) ap-east-1 : Asia Pacific (Hong Kong)
21) me-south-1 : Middle East (Bahrain)
22) af-south-1 : Africa (Cape Town)
(default is 3): 4

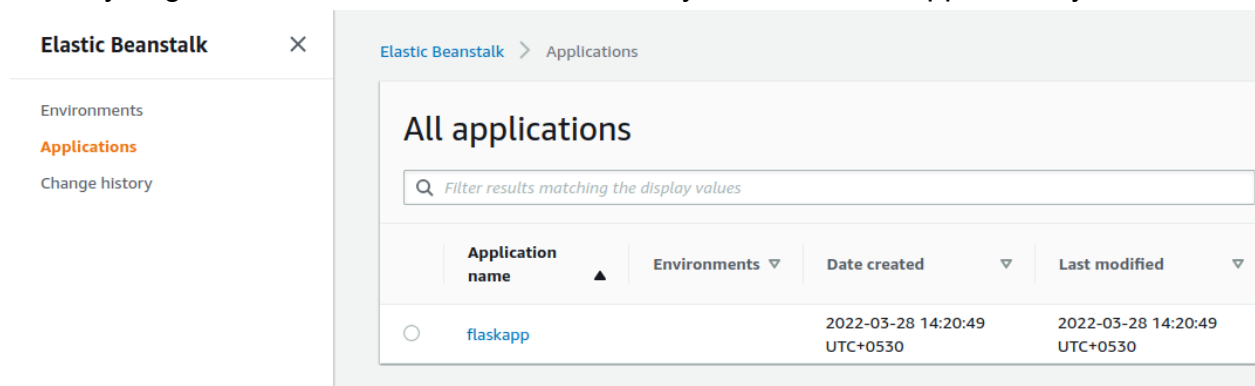
Enter Application Name
(default is "flaskapp"):
Application flaskapp has been created.

It appears you are using Python. Is this correct?
(Y/n): y
Select a platform branch.
1) Python 3.8 running on 64bit Amazon Linux 2
2) Python 3.7 running on 64bit Amazon Linux 2
3) Python 3.6 running on 64bit Amazon Linux (Deprecated)
(default is 1): 1

Cannot setup CodeCommit because there is no Source Control setup, continuing with initialization
Do you want to set up SSH for your instances?
(Y/n): n

```

Now if you got elastic beanstalk in aws console you will see the application you created.



Now lets create an environment. In terminal type eb create.

**eb create**

I kept most of the configurations to default as in the image just set [spot fleet request](#) to N. And it will create a set of aws resources such as security groups, load balancers, ec2



instances,s3 bucket for keeping the code .ect for you. It will take a few minutes to create all the resources.

```
Enter Environment Name
(default is flaskapp-dev):
Enter DNS CNAME prefix
(default is flaskapp-dev):

Select a load balancer type
1) classic
2) application
3) network
(default is 2):

Would you like to enable Spot Fleet requests for this environment? (y/N): N
```

Once you get all setup go to aws console and navigate to aws beanstalk and see the application status.

The screenshot shows the AWS Elastic Beanstalk console. On the left is a navigation menu with options like Environments, Applications, and Change history. The main area displays details for the 'flaskapp-dev' environment. It includes a 'Health' section with a green checkmark and 'Ok' status, a 'Running version' section showing 'app-220328\_163359567017' with an 'Upload and deploy' button, and a 'Platform' section showing 'Python 3.8 running on 64bit Amazon Linux 2/3.3.11' with a 'Change' button. Below these is a 'Recent events' table with three entries, all of type 'INFO' and message 'Deleted log fragments for this environment.'

Time	Type	Details
2022-03-28 16:50:29 UTC+0530	INFO	Deleted log fragments for this environment.
2022-03-28 16:46:49 UTC+0530	INFO	Deleted log fragments for this environment.
2022-03-28 16:44:34 UTC+0530	INFO	Deleted log fragments for this environment.

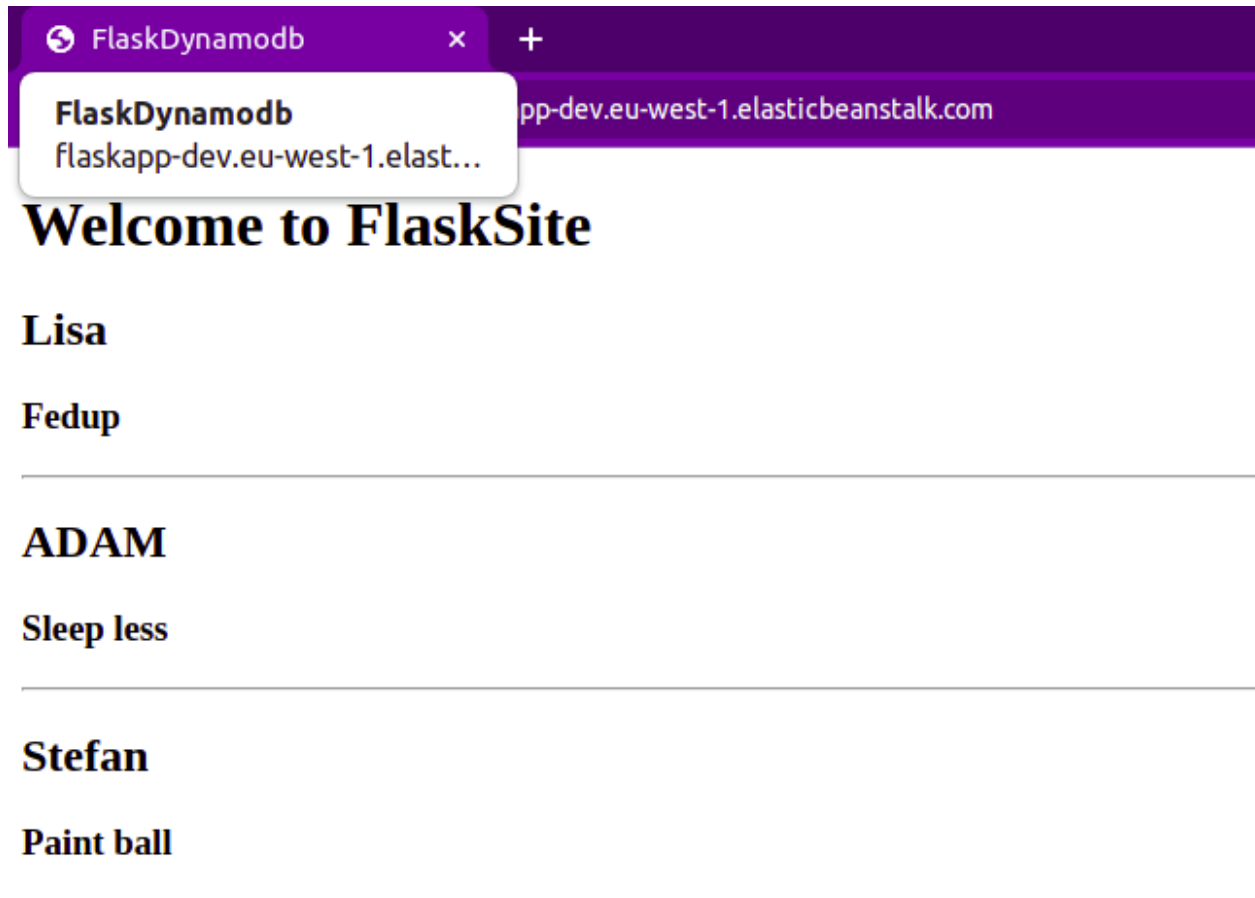
And also you can see the status via command `eb status` in terminal

**eb status**

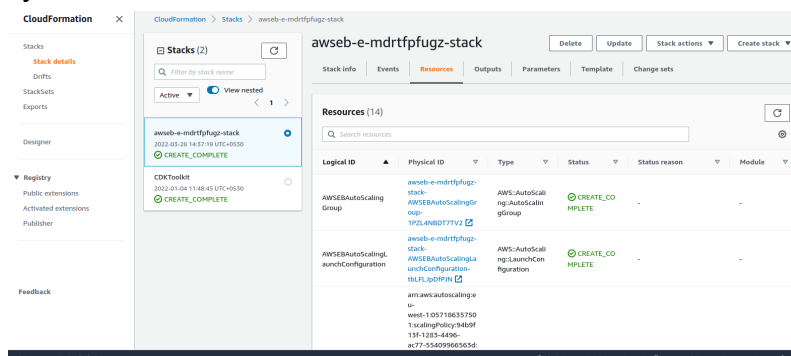
```
Environment details for: flaskapp-dev
Application name: flaskapp
Region: eu-west-1
Deployed Version: app-220328_163859169012
Environment ID: e-mdrtfufugz
Platform: arn:aws:elasticbeanstalk:eu-west-1::platform/Python 3.8
Tier: WebServer-Standard-1.0
CNAME: flaskapp-dev.eu-west-1.elasticbeanstalk.com
Updated: 2022-03-28 11:09:33.719000+00:00
Status: Ready
Health: Green
```

Now you can open the web page by typing `eb open`  
`eb open`

It will open a web browser with output like this.



If you navigate to cloudformation in aws console you will see all the resources created by the beanstalk.



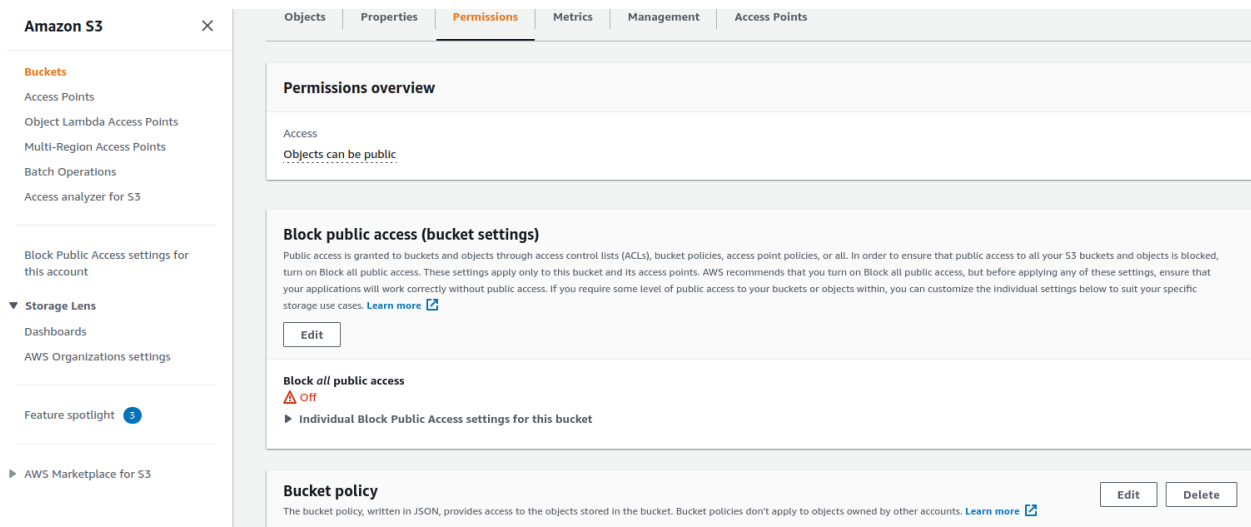
## Remove Deployment:

If you need to terminate the deployed service you can run below command

```
eb terminate --force
```

This will remove all the resources except s3 bucket. You have to remove the attached bucket policy and delete the bucket from aws console.

Navigate to s3 and there is a tab named permission in there you have the option to delete the bucket policy.



After that just delete the bucket.